



Data Wrangling

Exercises

1. Some sets have missing information for retail_price or pieces or both. This could be because the sets are free (giveaways), they aren't traditional lego sets (comic books, etc) or just because the information is missing. Filter the lego dataset based on the specifications below and save the result as lego using <-. Hence, you will overwrite the original lego object. In addition, describe the implications of removing these sets.

Your new lego tibble (data frame) should have:

- no missing pieces
- only contain sets with a nonzero number of pieces
- no missing retail_price
- only contain sets with a nonzero retail_price
- no missing year

```
25
26 # Exercise 1
27
28 ```{r}
29 lego_filtered <- lego %>%
30 filter(!is.na(pieces)) %>%
31 filter(pieces != 0) %>%
32 filter(!is.na(retail_price)) %>%
33 filter(retail_price != 0) %>%
34 filter(!is.na(year))
35
36 ```
37
```

31:24 | Chunk 3 | R Markdown

Console | Terminal | Background Jobs

R 4.3.1 · C:/Users/reemH/Downloads/DS/Labs/Lab2/lab-02-ReemAlsharabi/

Rows: 12214 Columns: 10

— Column specification —

Delimiter: ","

chr (6): id, name, themegroup, theme, subtheme, package

dbl (4): year, pieces, minifigs, retail_price

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
> lego_filtered <- lego %>%
+ filter(!is.na(pieces)) %>%
+ filter(pieces != 0) %>%
+ filter(!is.na(retail_price)) %>%
+ filter(retail_price != 0) %>%
+ filter(!is.na(year))
>
```



2. Arrange the dataset in descending order of *retail_price* and print the first three rows. Report in words the names of the three most expensive lego sets, their prices, and how many pieces each has.

```
37
38 # Exercise 2
39
40 ```{r}
41 lego_filtered %>%
42   arrange(desc(retail_price)) %>%
43   slice(1:3)%>%
44   print(width = Inf)
45 ```
46
47 Describe the three most expensive sets here.
48 the most expensive 3 lego sets are Millennium Falcon , Connections Kit and Death Star.
49
```

49:1 # Exercise 2 : R Markdown

Console Terminal Background Jobs

R 4.3.1 · C:/Users/reemH/Downloads/DS/Labs/Lab2/lab-02-ReemAlsharabi/

```
> lego_filtered %>%
+ arrange(desc(retail_price)) %>%
+ slice(1:3)%>%
+ print(width = Inf)
# A tibble: 3 × 10
  id      name      themegroup theme      subtheme      year pieces
<chr>   <chr>      <chr>      <chr>      <chr>      <dbl> <dbl>
1 75192-1 Millennium Falcon Licensed Star Wars Ultimate Collector Series 2017 7541
2 2000431-1 Connections Kit Educational Serious Play NA 2013 2455
3 75159-1 Death Star Licensed Star Wars Ultimate Collector Series 2016 4016
  minifigs package retail_price
<dbl> <chr> <dbl>
1      8 Box      800.
2      0 NA      755.
3     27 Box      500.
>
```

- Set: Millennium Falcon
Price: \$800
Pieces: 7,541
- Set: Connections Kit
Price: \$755
Pieces: 2,455
- Set: Death Star
Price: \$500
Pieces: 4,016



- It appears that the most expensive sets generally have more pieces. Use `mutate()` to create a new variable `price_per_piece`, representing the price in dollars per piece for each of the sets. Save the result as `lego`. Hence, you will overwrite the current `lego` object.

```
50
51 # Exercise 3
52
53 {r}
54 mutate( lego_filtered, "price_per_piece")
55 lego <- lego %>%
56   mutate(price_per_piece = retail_price / pieces)
57 print(lego, width = Inf)
58
59
```

57:25 Chunk 5 R Markdown

Console Terminal Background Jobs

R 4.3.1 · C:/Users/reemH/Downloads/DS/Labs/Lab2/lab-02-ReemAlsharabi/

	year	pieces	minifigs	package	retail_price	price_per_piece	
	<dbl>	<dbl>	<dbl>	<chr>	<dbl>	<dbl>	
1	2019	NA	0	Blister pack	NA	NA	
2	2019	2569	6	Box	200.	0.0778	
3	2019	1471	0	Box	150.	0.102	
4	2019	69	0	Box	9.99	0.145	
5	2019	178	0	Box	35.0	0.197	
6	2019	230	0	Box	50.0	0.217	
7	2019	98	0	Box	20.0	0.204	
8	2019	11	0	Box	9.99	0.908	
9	2019	15	0	Box	15.0	0.999	
10	2019	NA	0	NA	20.0	NA	

i 12,204 more rows
i Use `print(n = ...)` to see more rows
> |



4. Arrange the *lego* dataset in descending order of *price_per_piece* and return only the columns *name*, *themegroup*, *theme*, *pieces*, *price_per_piece*, and the first five rows. What do you notice about these sets?

```
61 {r}
62 # Filter out sets with low piece counts
63 lego_filtered <- lego %>% filter(pieces > 1)
64
65 # Create the price_per_piece variable
66 lego_filtered <- lego_filtered %>%
67   mutate(price_per_piece = retail_price / pieces)
68
69 # Arrange the dataset in descending order of price_per_piece and select specific columns
70 lego_sorted <- lego_filtered %>%
71   arrange(desc(price_per_piece)) %>%
72   select(name, themegroup, theme, pieces, price_per_piece)
73
74 # Print the first five rows
75 head(lego_sorted, 5)
76
77
```

75:21 | Chunk 6 | R Markdown

Console | Terminal | Background Jobs

R 4.3.1 · C:/Users/reemH/Downloads/DS/Labs/Lab2/lab-02-ReemAlsharabi/

```
> # Arrange the dataset in descending order of price_per_piece and select specific columns
> lego_sorted <- lego_filtered %>%
+   arrange(desc(price_per_piece)) %>%
+   select(name, themegroup, theme, pieces, price_per_piece)
>
> # Print the first five rows
> head(lego_sorted, 5)
# A tibble: 5 × 5
  name                themegroup theme    pieces price_per_piece
<chr>                <chr>    <chr>    <dbl>    <dbl>
1 Adventurers combined set Educational Dacta      3      29.8
2 Collectable Display Set 2 Licensed   Star Wars    2      25.0
3 Building Table          Basic      Basic      3      16.5
4 Infrared Transmitter with USB Cable Educational Education  2       15
5 Play Table              Basic      Basic      3       15
>
```

- The sets generally have a small number of pieces, ranging from 2 to 3.
- Despite their small piece counts, these sets have a higher price per piece, ranging from \$15 to \$29.8.
- The sets cover diverse themes, including Educational, Licensed (Star Wars), Basic, and Education.
- These sets may possess unique features, specialized designs, educational value, or collectible qualities that justify their higher price per piece.

Overall, these observations suggest that these sets with small piece counts and higher prices per piece offer distinctive qualities or cater to specific purposes, making them stand out within the LEGO product range.



5. Return a tibble containing the cheapest and most expensive lego sets (based on *retail_price*) in each subtheme, considering only sets with the Lord of the Rings theme.

```
82
83 {r}
84 library(dplyr)
85
86 # Filter dataset for "The Lord of the Rings" theme and non-missing retail prices
87 lego_filtered_lotr <- lego[lego$theme == "The Lord of the Rings" &
88   !is.na(lego$retail_price), ]
89
90 # Group the filtered dataset by subtheme
91 grouped_data <- lego_filtered_lotr %>%
92   group_by(subtheme)
93
94 # Calculate the cheapest and most expensive sets in each subtheme
95 result <- grouped_data %>%
96   summarize(
97     cheapest_set = min(retail_price),
98     most_expensive_set = max(retail_price)
99   )
100 result_tibble <- as_tibble(result)
101 result_tibble
```

100:1 | Chunk 7 | R Markdown

Console | Terminal | Background Jobs

R 4.3.1 · C:/Users/reemH/Downloads/DS/Labs/Lab2/lab-02-ReemAlsharabi/

```
+ most_expensive_set = max(retail_price)
+ )
>
> # Convert the result to a tibble
> result_tibble <- as_tibble(result)
>
> # View the tibble
> result_tibble
# A tibble: 4 × 3
  subtheme      cheapest_set most_expensive_set
  <chr>          <dbl>          <dbl>
1 The Fellowship of the Ring      0          80.0
2 The Return of the King      20.0         100.
3 The Two Towers           4         200.
4 NA                NA                NA
> |
```

id	name	themegroup	theme	subtheme	year	pieces	minifigs	package	retail_price	price_per_piece
1	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
2	10237-1	Tower of Orthanc	Licensed	The Lord of the Rings The Two Towers	2013	2359	5	Box	199.99	0.08477745
3	79005-1	The Wizard Battle	Licensed	The Lord of the Rings The Fellowship of the Ring	2013	113	2	Box	12.99	0.11495575
4	79006-1	The Council of Elrond	Licensed	The Lord of the Rings The Fellowship of the Ring	2013	243	4	Box	29.99	0.12341564
5	79007-1	Battle at the Black Gate	Licensed	The Lord of the Rings The Return of the King	2013	656	5	Box	59.99	0.09144817
6	79008-1	Pirate Ship Ambush	Licensed	The Lord of the Rings The Return of the King	2013	756	9	Box	99.99	0.13226190
7	30210-1	Frodo with cooking corner	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	33	1	Polybag	3.99	0.12090909
8	30211-1	Uruk-Hai with ballista	Licensed	The Lord of the Rings The Two Towers	2012	21	1	Polybag	4.00	0.19047619
9	5000202-1	Elrond	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	6	1	Polybag	0.00	0.00000000
10	9469-1	Gandalf Arrives	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	83	2	Box	12.99	0.15650602
11	9470-1	Shelob Attacks	Licensed	The Lord of the Rings The Return of the King	2012	227	3	Box	19.99	0.08806167
12	9471-1	Uruk-Hai Army	Licensed	The Lord of the Rings The Two Towers	2012	257	6	Box	29.99	0.11669261
13	9472-1	Attack On Weathertop	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	430	5	Box	59.99	0.13951163
14	9473-1	The Mines of Moria	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	776	9	Box	79.99	0.10307990
15	9474-1	The Battle of Helm's Deep	Licensed	The Lord of the Rings The Two Towers	2012	1368	8	Box	129.99	0.09502193
16	9476-1	The Orc Forge	Licensed	The Lord of the Rings The Fellowship of the Ring	2012	366	4	Box	39.99	0.10926230



6. Use `group_by()` and `summarize()` to create a new tibble with one row for each year, and columns for the year, the number of sets released in that year, and the median price per piece for sets from that year. Save this resulting tibble as an object named `yearly_trends`.

```
102
103 # Exercise 6
104
105 # Remove rows with missing or zero values in retail_price or pieces columns
106 lego_clean <- lego %>%
107   filter(!is.na(retail_price) & !is.na(pieces) & retail_price != 0 & pieces != 0)
108
109 # Recalculate yearly trends with the cleaned dataset
110 yearly_trends <- lego_clean %>%
111   group_by(year) %>%
112   summarize(
113     num_sets = n(),
114     median_price_per_piece = median(retail_price / pieces)
115   )
116
117 # Convert the result to a tibble
118 yearly_trends <- as_tibble(yearly_trends)
119
120 # View the yearly_trends tibble
121 yearly_trends
122
123
124
```

115:12 Chunk 8

R 4.3.1 - C:/Users/reemh/Downloads/DS/Lab2/lab-02-ReemAlsharabi/

```
> yearly_trends
# A tibble: 58 x 3
  year num_sets median_price_per_piece
  <dbl> <int> <dbl>
1 1961      2 0.0340
2 1963      1 0.0308
3 1964      7 0.0113
4 1965      2 0.883
5 1966     18 0.0503
6 1967      8 0.0356
7 1968      3 0.0443
8 1969      3 0.0443
9 1970      4 0.0660
10 1971     13 0.0308
```

	year	num_sets	median_price_per_piece
1	1961	2	0.03402778
2	1963	1	0.03076923
3	1964	7	0.01126437
4	1965	2	0.88255814
5	1966	18	0.05031745
6	1967	8	0.03556094
7	1968	3	0.04432373
8	1969	3	0.04432373
9	1970	4	0.06604144
10	1971	13	0.03076923
11	1973	3	0.02448454
12	1974	3	0.06410256
13	1975	1	0.71190476
14	1976	5	0.07425743
15	1977	4	0.57083333

Showing 1 to 16 of 58 entries, 3 total columns



7. Create a plot of the median price per piece over time using the *yearly_trends* tibble. Size points according to the number of sets produced in that year. Adjust transparency, color, etc as appropriate and remember the principles of effective data visualization. Comment on what you observe.
- From 1960 to 1980: The median price per piece was consistently low, below 0.25, with a few high outliers indicating occasional sets with higher prices per piece.
 - From 1980 to 1990: The median price per piece showed an increase and approached the 0.25 threshold, suggesting that the average price per piece became more aligned with this value during this decade.
 - From 1990 to 2020: The median price per piece remained relatively low and did not exceed 0.25. There were no outliers during this period, indicating a consistent price per piece for LEGO sets without extremely high or low values.


These findings suggest that LEGO sets have generally maintained an affordable price per piece over the years, with some fluctuations and a trend toward closer alignment with the 0.25 price point.

```
{r}
library(ggplot2)

# Filter out rows with missing values in the yearly_trends tibble
yearly_trends_filtered <- yearly_trends %>%
  filter(!is.na(year) & !is.na(median_price_per_piece))

# Plotting the median price per piece over time
ggplot(yearly_trends_filtered, aes(x = year, y = median_price_per_piece)) +
  geom_point(aes(size = num_sets), alpha = 0.7, color = "steelblue") +
  scale_size(range = c(3, 10)) +
  labs(x = "Year", y = "Median Price per Piece", size = "Number of Sets") +
  theme_minimal()
```

Comment on what you observe in the plot above.
The analysis of the median price per piece over time reveals that from 1960 to 1980, the median price per piece was consistently low, with occasional high outliers. However, from 1980 to 1990, there was a shift towards prices closer to the 0.25 threshold. From 1990 to 2020, the median price per piece remained consistently low, without outliers. These findings suggest that LEGO sets have generally maintained an affordable price per piece over the years, with some fluctuations and a trend towards closer alignment with the 0.25 price point.



Submission

Knit to PDF to create a PDF document. Stage and commit all remaining changes, and push your work to GitHub. Make sure all files are updated on your GitHub repo. Only upload your PDF document to Blackboard. Before you submit the uploaded document, mark where each answer is to the exercises.